



## Introduction to Python for data scientists

workshop, 20.08.2019

# Slides and files

<https://drive.google.com/drive/folders/1xCaSxCax04hig8uh47prJpGpgCgn1eT-?usp=sharing>

# Goals

- Have a working Python environment set up
- Be able to run Python code
- Know basic syntax / know where to look for help
- Be able to install and use new packages
- Have Jupyter set up
- Load, save and manipulate tabular data in Python
- Calculate using data, plot results

# Agenda

10 min **Introduction**

80 min **Basic Python syntax with exercises**

10 min **Break**

10 min **Modules & packages**

20 min **Jupyter**


60 min **Introduction to data analysis with Pandas**

10 min **Break**

Remaining time: **Final exercise, questions**

# Introduction

# Data science?



**mobile apps, websites, business software, robots,**  
**autonomous cars, crypto currencies, satellites, nuclear plants**  
 teams of professional developers  
 thousands of lines of code  
 project management, release dates, leanagilescrum  
 requirements  
 software design, architecture, patterns, styles  
 clients, users  
 result: tested, working, re-usable, safe code

**data analysis, statistics, simulations,**  
**physical, mathematical computations**  
 scientist(s), co-authors  
 as few lines of code as possible  
 project management !?  
 problems, ideas, data  
 exploration  
 scientific community  
 result: (reproducible) findings

# Data science



easy to learn

quick results

reliable results

easy to share

reproducible

data science

**data analysis, statistics, simulations,  
physical, mathematical computations**  
scientist(s), co-authors  
as few lines of code as possible  
project management !?  
problems, ideas, data  
exploration  
scientific community  
result: (reproducible) findings

# What is Python?

- “Python is an interpreted, object-oriented, high-level programming language with dynamic semantics.”
- Invented in the 90s by Guido van Rossum (NL)
- Free / Open Source
- Has become de-facto standard for data science (along with other languages and tools...)






# Python is simple

- Easy to learn, quick to get results with
- Intuitive syntax, readable code
- Hard to break
- Runs on every major platform (windows, macOS, linux etc.)
- Huge offer of packages (“add-ons”) to choose from (most problems are already solved)
- Well-tried in professional software development
- Large active community, many tutorials etc.


# Python is trendy

## Growth of major programming languages

Based on Stack Overflow question views in World Bank high-income countries



# Everybody is using Python



**moz://a**

**NETFLIX**



**NOKIA**

*Disney*



**Uber**

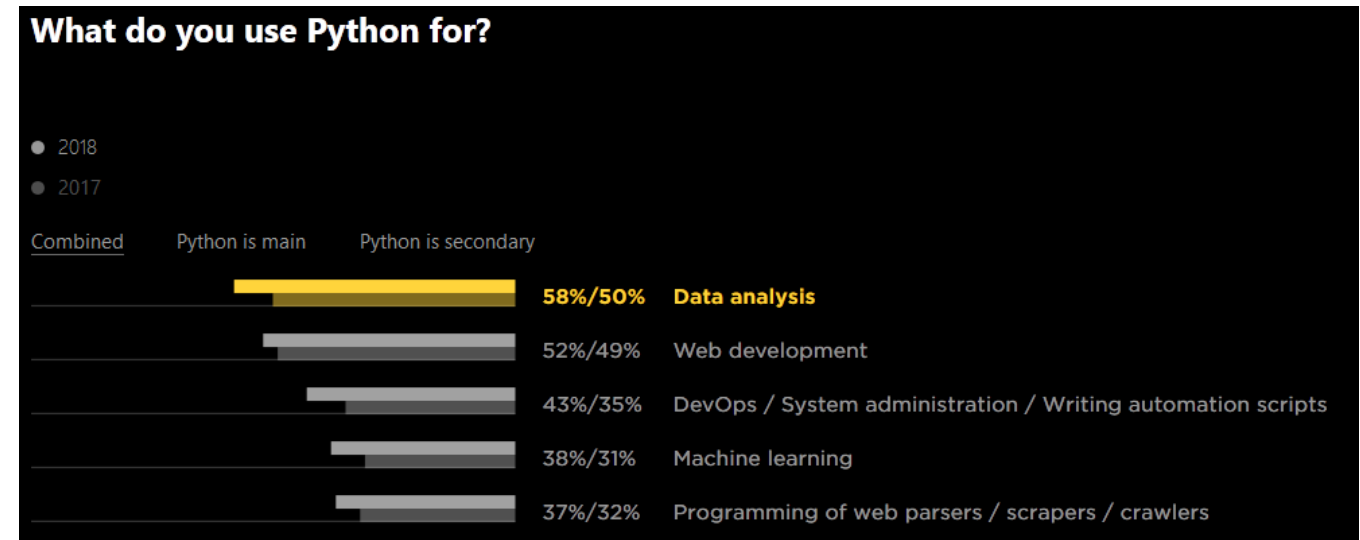


# Python: de-facto standard for data analysis

KDnuggets Software Poll (n>1,800)

Platform	2019 % share	2018 % share	% change
<b>Python</b>	<b>65.8%</b>	<b>65.6%</b>	<b>0.2%</b>
R Language	46.6%	48.5%	-4.0%
SQL Language	32.8%	39.6%	-17.2%
Java	12.4%	15.1%	-17.7%
Unix shell/awk	7.9%	9.2%	-13.4%
C/C++	7.1%	6.8%	3.7%
Javascript	6.8%	na	na
Other programming and data languages	5.7%	6.9%	-17.1%
Scala	3.5%	5.9%	-41.0%
Julia	1.7%	0.7%	150.4%
Perl	1.3%	1.0%	25.2%

<https://www.kdnuggets.com/2019/05/poll-top-data-science-machine-learning-platforms.html>



<https://www.jetbrains.com/research/python-developers-survey-2018/>

# Running python

1. **Interactive mode:** Running code in the console. On Windows:  
⌘ win + R → “cmd” → `ipython`
2. **Scripts:** Executing a python-script (filenames with .py extension) → console: `python filename.py`
3. **Notebooks:** Running code in a “notebook”, e.g. Jupyter

*Exercise (2 mins):*

Start Python in interactive mode

# Python syntax

# Python syntax

- Variables: `x = 1 + 2`
- Basic blocks: `=`, `==`, `>`, `<`, `>=`, `<=`, `!=`, `not`, `and`, `or`, `(`, `)`
- Each line is a command
- `# comments` will not be executed
- `x = 1 + 2 # anything after “#” is a comment`
- Python3 supports Unicode: `肉 = 1 + 2`
- use `print(...)` to output expressions

# Python syntax: Basic types & operations

- None type
- Booleans: `True`, `False`
- Numbers: `int` (`-1`, `0`, `1`, ...) `float` (`-1.1283`, ...) , `complex` (...)
- Text: `str` (`"this is a text"`, `'this is a text'`)

- Sequences / sets:

list: `[0, 1, 2]`, `['a', 'b', 1]`, `[True, False]`, `[]`

tuple: `(0, 1, 2)`, `('a', 'b', 1)`

set: `{0, 1, 2}`, `{True, 'a', 1}`, ~~`{'a', 1, 'a'}`~~

- Maps:

map { `'a': 0.284`, `'b': True`, `1: "xy"` }



# Python syntax: Booleans



- `True, False` *# case sensitive!*
- `x = True`
- `y = False`
- `x or y` *# True*
- `x and y` *# False*
- `not x` *# False*
- `not y` *# True*
- `not ( x and y )` *# True*

# Python syntax: Numbers

Operation	Result
<code>x + y</code>	sum of x and y
<code>x - y</code>	difference of x and y
<code>x * y</code>	product of x and y
<code>x / y</code>	quotient of x and y
<code>x % y</code>	remainder of x / y
<code>-x</code>	x negated
<code>+x</code>	x unchanged
<code>abs(x)</code>	absolute value or magnitude of x
<code>int(x)</code>	x converted to integer
<code>float(x)</code>	x converted to floating point
<code>pow(x, y)</code>	x to the power y
<code>x ** y</code>	x to the power y

# Python syntax: Numbers

*Exercise (5 mins):*



DIN 19584-1:2012-10

What is the area (in  $\text{m}^2$ ) of this DIN 19584 – A – D 400 manhole cover?

*Hints:*  $\pi \approx 3.14$ ,  $r = 0.5 * 785 \text{ mm}$

# Python syntax: strings

- `x = "FG INNO"`
- `len(x) # 7`
- `"I" in x # True`
- `x.count("N") # 2`
- `x[1] # G`
- `x = x + " is great" # x = "FG INNO is great"`
- `y = "!!!"`
- `x + y # "FG INNO is great!!!"`
- "You can insert new lines `\n` and tabs `\t` etc."

# Python syntax: lists

- `x = [1, 5, 7, 3, 5, 9]`
- `len(x) # 6`
- `sum(x) # 30`
- `max(x) # 9`
- `3 in x # True`
- `x.count(5) # 2`
- `x.sort() # x = [1, 3, 5, 5, 7, 9]`
- `x + [10, 11] # [1, 3, 5, 5, 7, 9, 10, 11]`

# Python syntax: tuples & sets

## Quiz

What is

(a) `len( set( [1,1,2,2,3,3] ) )`

(b) `(2,1,3).sort()`

?

# Python syntax: list slicing

- `list[start:stop:step]`
- `x = ["DIN", "ISO", "DIN", "ETSI", "IETF"]`
- `x[1:3] # ["ISO", "DIN", "ETSI"]`
- `x[3:] # ["ETSI", "IETF"]`
- `x[:1] # ["DIN", "ISO"]`
- `x[:-3] # ["DIN", "ISO"]`
- `x[1:4:2] # ["ISO", "ETSI"]`

# Python syntax: list comprehension

- Create a list by giving a functional definition of the elements
- `x = [i for i in range(1,5)] # [1,2,3,4]`
- `y = [a/2 for a in range(0,3)] # [0, 0.5, 1]`
- `z = [(n, 2*n) for n in y] # [(0,0), (0.5, 1), (1,2)]`
- `z_1 = [(n, 2*n) for n in y if n<1] # [(0,0), (0.5, 1)]`
- `z_2 = [(n, 2*n) for n in y][::2] # [(0,0), (1, 2)]`



# Python syntax: lists

*Exercise (10 mins):*

What is the sum of every second power of two:  
 $2^n$  for  $n = 0..10$  ?

# Python syntax: maps

- `map = { key1: value1, key2: value2}` *# keys are unique*
- `map[key]` *# value*
- `levels = { "DIN": "DE", "DKE": "DE", "CEN": "EU", "CENELEC": "EU"}`
- `level["DIN"]` *# DE*
- `level["ETSI"] = "EU"`
- `level.keys()` *# DIN, DKE, CEN, CENELEC*
- `level.values()` *# DE, DE, EU, EU*
- `x = {i: i**2 for i in range(2,6)}` *# {2:4, 3:9, 4:16, 5:25}*

# Python syntax: control flows

- The usual suspects: `if`, `else`, `while`, `for`
- Inside these constructs, statements are intended (e.g. with a tab):

```
if x > 5:  
    print("x is too big")  
    x = 5  
else:  
    print("x is ok")
```

# Python syntax: if, else, elif

```
if x > 5:
    print("x is too big")
elif x < 2:
    print("x is too small")
else:
    print("x is ok")

s = "negative" if x < 0 else "positive"
```

# Python syntax: while

- Execute code as long as an expression is **True**

```
x = 0
while x < 10:
    print(x)
    x = x + 1
```

```
# 0 1 2 3 4 5 6 7 8 9
```

# Python syntax: for

- Same as in list comprehensions, used to iterate over a **sequence of items**

```
for x in range(10):  
    print(x)
```

```
for x in [i for i in range(10) if i % 2 == 0]:  
    print(x)
```

```
for w in "an example sentence".split(" "):  
    print(w)
```

# Python syntax: control flow

*Exercise (15 mins):*

How many prime numbers are  $< 100$ ?

# Python syntax: functions

- Re-usable code blocks with parameters

*# a function that returns all prime numbers < limit*

```
def primes(limit):  
    p = [n for n in range(2, limit + 1)  
         if not any([n % x == 0 for x in range(2, n)])]  
  
    return p
```

```
print(primes(1000))  
print(373 in primes(400))
```



# Python syntax: functions

```
def function_name(x, y, z):
```

```
    ...
```

```
# parameters be optional and can have default values
```

```
def f(x, y, z=5):
```

```
    return x+y+z
```

```
f(1,2) # 8
```

```
f(1,2,3) # 6
```

```
f(y=2, x=1, z=3) # 6
```

# Python syntax: (parameter) unpacking

```
a, b, c, d = [1, 2, 3, 4] # a=1, b=2, ...
```

```
def f(x, y, z):  
    return [x+y, z]
```

```
a, b = f(1,2,3) # a=3, b=3
```

```
f(*[1,2,3]) == f(1,2,3) # True, pass parameters as a list
```

```
f(*(1,2,3)) # ... a tuple
```

```
f(**{'x': 1, 'z': 3, 'y': 2}) # ... or named as a map
```

# Python syntax: (parameter) unpacking

```
import math
```

```
# example: transformation from cartesian to spherical coordinates
```

```
def spherical(x, y, z):  
    r = math.sqrt(x**2 + y**2 + z**2)  
    phi = math.arctan(y/x)  
    theta = math.arccos(z/r)  
  
    return (r, theta, phi)
```

```
# some_object.coordinates() = (x, y, z)
```

```
r, phi, theta = spherical(*some_object.coordinates())
```

# Python syntax: functions

*Exercise (10 mins):*

Write a function that transforms spherical coordinates to cartesian coordinates:  $(r, \theta, \phi) \rightarrow (x, y, z)$

*Hint:*

$$\begin{aligned}x &= r \sin \theta \cos \varphi \\y &= r \sin \theta \sin \varphi \\z &= r \cos \theta\end{aligned}$$

Use `import math` to load the `math` module, and the `math.sin(x)` and `math.cos(x)` functions.

# Python syntax: functions

## Quiz

How would you output just the resulting `y` of  
`cartesian(*(coords))` ?

# Python syntax: lambda expressions

```
def plus(x, y):  
    return x + y  
plus(1,2) # 3
```

```
minus = lambda x,y: x - y  
minus(3,1) # 2
```

```
map( lambda x: x**2 + (x/2), [1,2,3,4] )  
# [1.5, 5.0, 10.5, 18.0]
```

**10 min break**

# **modules & packages**



# Python syntax: Modules

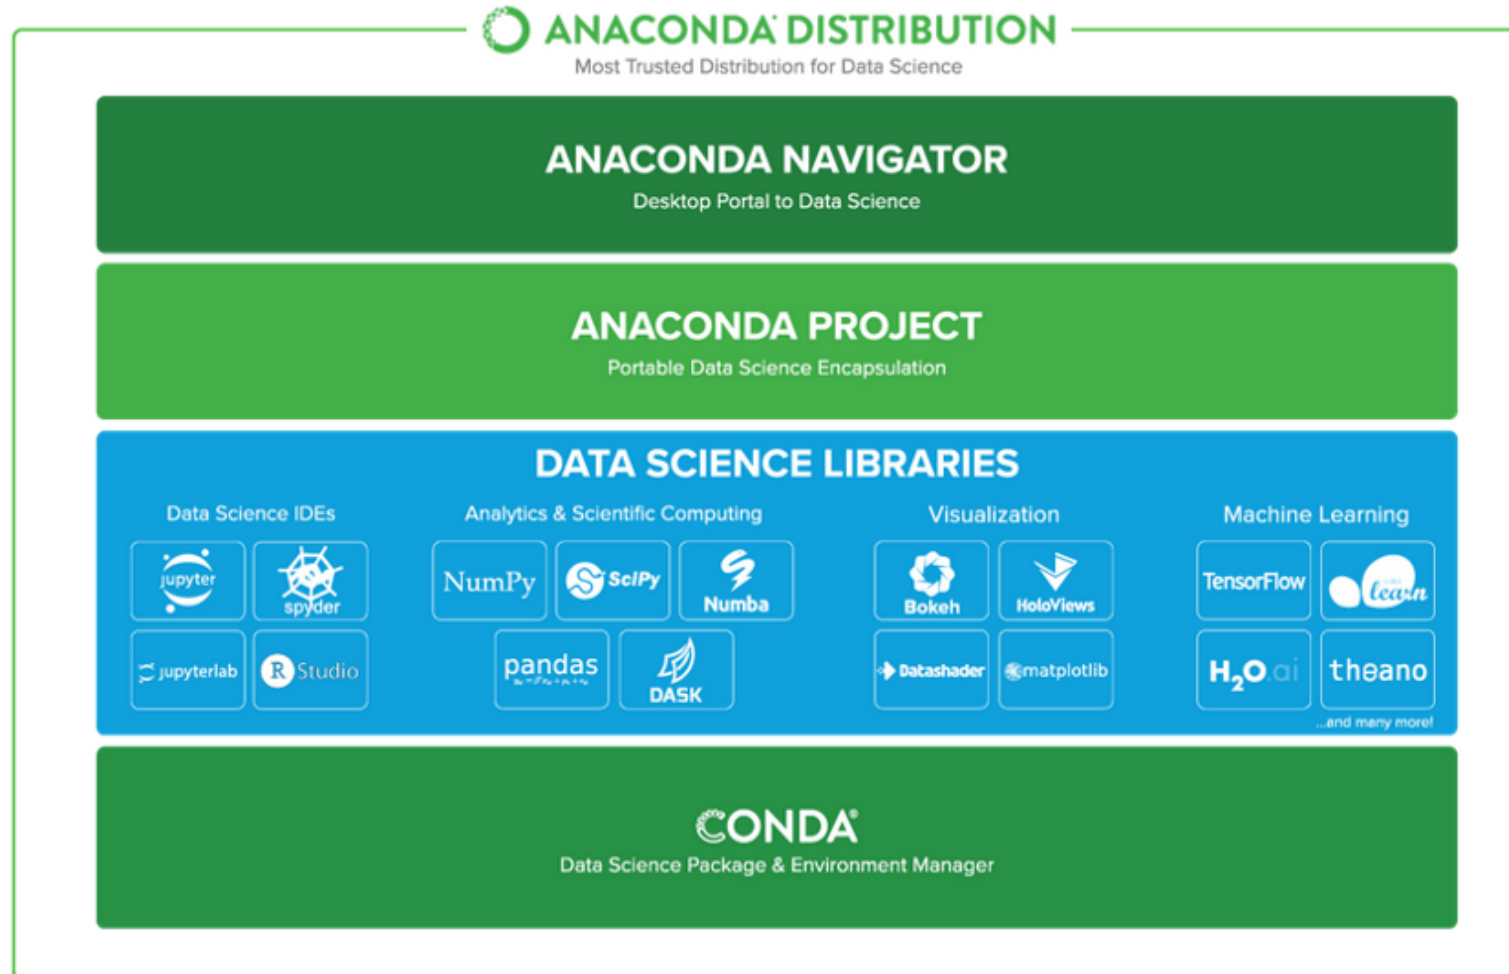
```
import math
import math as m
from math import cos
from math import sin, pi
from math import pi as  $\pi$ 
```

```
math.sin(0)
m.sin(0)
cos(0)
sin(0)
print(pi)
print( $\pi$ )
```

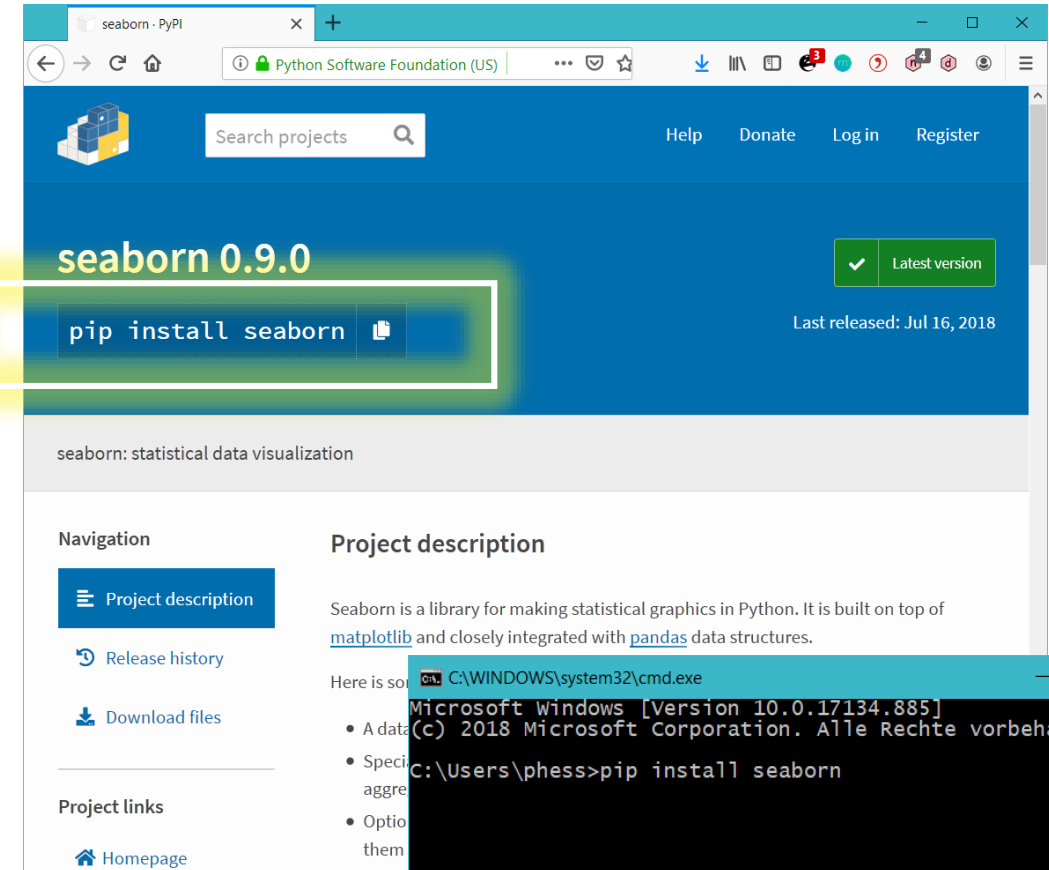

# Modules & packages

- Default **modules** that are always included in Python: <https://docs.python.org/3/library/index.html>. Examples:
  - math – mathematical functions
  - time – time access and conversions
  - random – generate pseudo-random numbers
- **Packages** (contain modules), open source, listed on <https://pypi.org>. Currently ~190,000 projects. Examples:
  - Numpy – (fast) scientific computing with Python
  - Pandas – data analysis library
  - TensorFlow – machine learning
  - Scikit-Learn – machine learning
- If you are using the Anaconda distribution, a lot of packages will already be installed!

# Modules & packages



# Modules & packages



# Modules & packages

```
import seaborn as sns
```

```
# some magic so the plot will be shown when produced in the console  
%matplotlib qt
```



```
# Load an example dataset with long-form data  
d = sns.load_dataset("fmri")
```

```
# Plot  
sns.lineplot(x="timepoint", y="signal",  
             hue="region", style="event", data=d)
```

# Jupyter

# Jupyter

- Write Python code in a “notebook” in your browser
- Output is included in the notebook
- Easy to share notebooks



# Jupyter notebooks

[jupyter\\_examples.ipynb](#)



# Running python

1. **Interactive mode:** Running code in the console. On Windows:  
⌘ win + R → "cmd" → `ipython`
2. **Scripts:** Executing a python-script (filenames with .py extension) → console: `python filename.py`
3. **Notebooks:** Running code in a "notebook", e.g. Jupyter

*Exercise (2 mins):*

Start Jupyter: ⌘ win → "Jupyter" → enter

# Jupyter notebooks


*Exercise (20 mins):*

- Create a new Jupyter notebook called „<your name>“ on your desktop.
- Copy and run one of the plotting examples from <https://seaborn.pydata.org/examples/index.html>
- Upload your notebook to the cloud at <https://tubcloud.tu-berlin.de/s/yXken8P3toea5XR>
- Look at and run some of the others' examples

# **Introduction to data analysis with Pandas**

# Pandas

- Part of the Anaconda distribution
- The “Excel”, “R data.tables” of Python
- Easy loading / saving / manipulation of tabular data
- Integrated easy plotting
- Querying data
- Support for time series
- ...



# Pandas: DataFrame, Series

- Core classes: **DataFrame** and **Series**
- **Series**: An ordered, indexed sequence of arbitrary values (numbers, strings, dates, ....)
- **DataFrame**: A two dimensional datatype with rows, columns and an index – a table. Or: a collection of Series (columns) with the same index

Series

	apples
0	3
1	2
2	0
3	1

+

Series

	oranges
0	0
1	3
2	7
3	2

=

DataFrame

	apples	oranges
0	3	0
1	2	3
2	0	7
3	1	2

# Pandas

- **Creating and exploring a dataset**
- **Indexing:** Selecting a certain data item from a Series, or a data item or sequence of data items from a DataFrame
- **Aggregation:** Running functions on sequences of data items
- **Grouping**
- **Selection:** Selecting items/sequences using logic
- **Plotting**

# Pandas

[pandas\\_intro.ipynb](#)

# Pandas: Series

*Exercise (10 mins):*

Generate a random Series that is correlated ( $\geq 0.5$ ) to the prime numbers below 100



# Pandas: indexing Series

- **Series** → Value (1d)

```
s1 = Series(range(5))
```

```
s1[2] # 2
```

```
s2 = Series(  
    range(5),  
    index=['A', 'B', 'C', 'D', 'E']  
)
```

```
s2['D'] # 3
```

<i>index</i>	
0	0
1	1
2	2
3	3
4	4

<i>index</i>	
A	0
B	1
C	2
D	3
E	4

# Pandas: indexing DataFrames

- **DataFrame** → Series (2d) / Value (1d)

```
df = DataFrame({
    'a': range(5),
    'b': range(2,7)
})
```

`df['b']` # 2 3 4 5 6 (*column*)

`df.iloc[0]` # 0 2 (*row*)

`df.iloc[1,1]` # 3 (*value*)

`df.iloc[::2]` # *every second row*

<i>index</i>	<b>a</b>	<b>b</b>
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

<i>index</i>	<b>a</b>	<b>b</b>
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

<i>index</i>	<b>a</b>	<b>b</b>
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

<i>index</i>	<b>a</b>	<b>b</b>
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

# Pandas: indexing DataFrames

- **names or numbers**

```
df1 = DataFrame([
    range(5),
    range(2,7)
])

df2 = DataFrame(
    { 'a': range(5),
      'b': range(2,7),
      'c': range(3,8) },
    index=['A','B','C','D','E']
)
```

<i>index</i>	0	1
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

<i>index</i>	a	b	c
A	0	2	3
B	1	3	4
C	2	4	5
D	3	5	6
E	4	6	7

# Pandas: indexing DataFrames

- **names or numbers**

```
df1[0] # first column
df1.iloc[0] # first row
df1.iloc[0,1] # row=0, col=1
df1.iloc[::2,1] # every second item of col=1
```

<i>index</i>	0	1
0	0	2
1	1	3
2	2	4
3	3	5
4	4	6

```
df2['a']
df2[['a','c']]
df2.loc['A'] == df2.iloc[0] # True
df2.loc[['C','E'],['a','b']] # 2 4
                                # 4 6
```

<i>index</i>	a	b	c
A	0	2	3
B	1	3	4
C	2	4	5
D	3	5	6
E	4	6	7

# Pandas: aggregating DataFrames

**Series** → Value (1d): `Series(range(5)).sum()` # 10

**DataFrame** → Series (2d) / Value (1d)

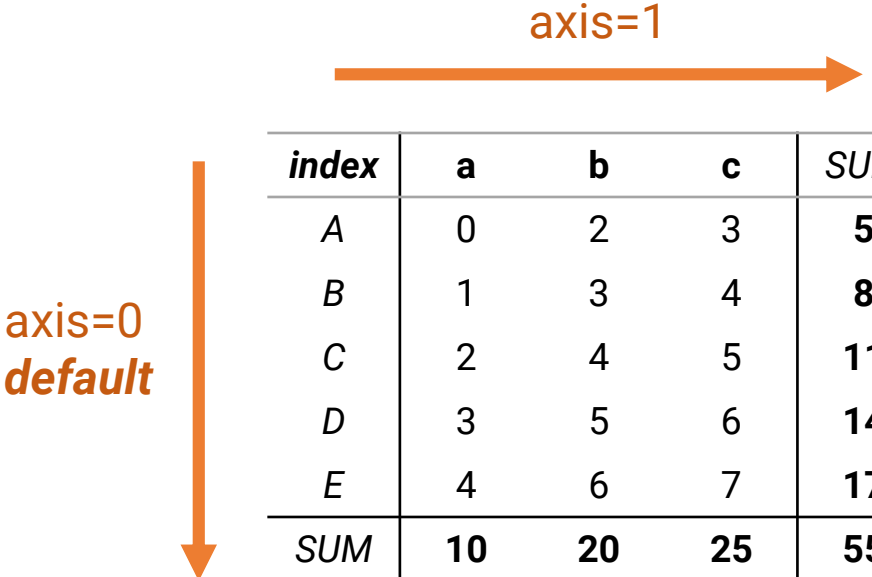
`df2['a'].sum()` # 10

`df2.sum()` # 10 20 25

`df2['A'].sum()` # 5

`df2.sum(axis=1)` # 5 8 11 14 17

`df2.sum().sum()` # 55



axis=1

<i>index</i>	<b>a</b>	<b>b</b>	<b>c</b>	<i>SUM</i>
A	0	2	3	<b>5</b>
B	1	3	4	<b>8</b>
C	2	4	5	<b>11</b>
D	3	5	6	<b>14</b>
E	4	6	7	<b>17</b>
<i>SUM</i>	<b>10</b>	<b>20</b>	<b>25</b>	<b>55</b>

axis=0  
default

**Other common aggregations:** mean, median, mode, max, min, var, std

<https://pandas.pydata.org/pandas-docs/stable/reference/frame.html#computations-descriptive-stats>

# Pandas: aggregating DataFrames

- **Compute any aggregation:**

- `df2.agg(numpy.sum) # == df2.sum()`
- `df2.agg(numpy.sum, axis=1) # == df2.sum(axis=1)`
- `df2.agg(numpy.log) # 5x3 df with log(item)`
- `df2.agg([numpy.sum, numpy.mean, numpy.std])`
- `df2.agg(lambda x: x**2 + x)`

# Pandas: aggregating DataFrames

- **Custom aggregations on axes**

*# aggregate each column with A - E*

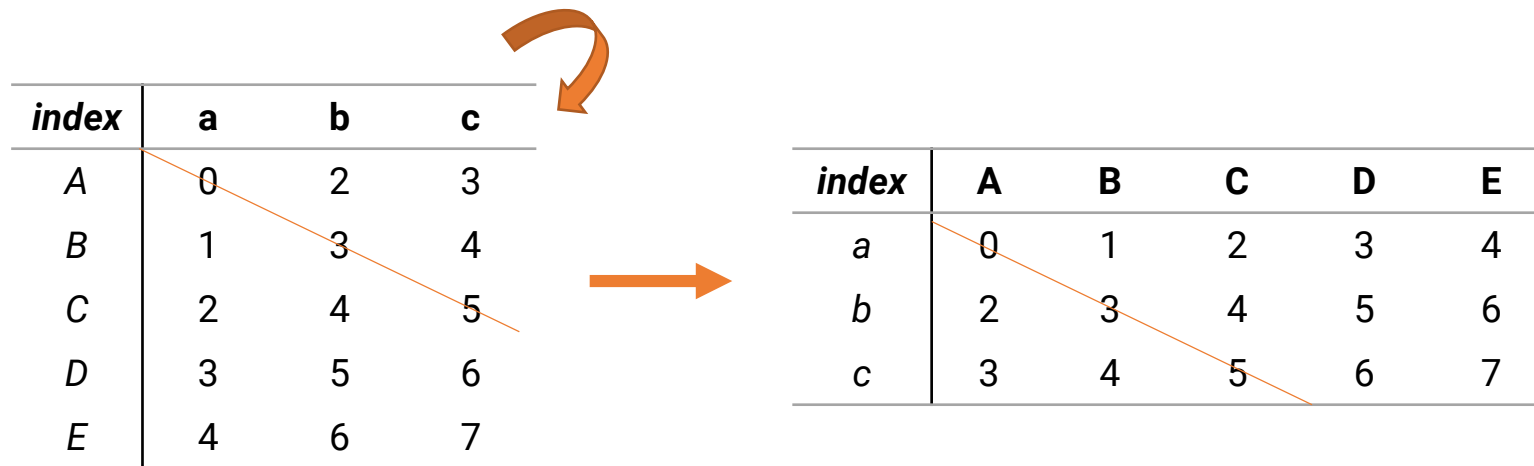
```
df2.apply(lambda rows: (rows['A'] - rows['E']))
```

*# aggregate each row with a / c*

```
df2.apply(lambda cols: (cols['a']/cols['c']), axis=1)
```

# Pandas: DataFrame transposition

- `df.T`



The diagram illustrates the transposition of a DataFrame. On the left, a DataFrame with index labels A, B, C, D, E and column labels a, b, c is shown. An orange arrow points from this DataFrame to the transposed DataFrame on the right. A curved orange arrow above the arrow indicates the transformation. The transposed DataFrame has index labels a, b, c and column labels A, B, C, D, E. The data values are swapped between rows and columns.

<i>index</i>	<b>a</b>	<b>b</b>	<b>c</b>
A	0	2	3
B	1	3	4
C	2	4	5
D	3	5	6
E	4	6	7

<i>index</i>	<b>A</b>	<b>B</b>	<b>C</b>	<b>D</b>	<b>E</b>
a	0	1	2	3	4
b	2	3	4	5	6
c	3	4	5	6	7



# Pandas: DataFrame indexing & aggregation

*Exercise (10 mins, choose one):*

- Get all 2019 innovation expenditures of sectors with numbers starting with “C”
- Get the innovation expenditures of “C 26” from 2010 to 2019
- Imagine that data was corrupted in every 2<sup>nd</sup> year. Use only the 1<sup>st</sup>, 3<sup>rd</sup>, ... year to calculate mean expenditures.
- Calculate the average change in expenditures since 2006 for sectors “C ..”

# Pandas: grouping

dfg =

<i>index</i>	<b>gender</b>	<b>age</b>	<b>height</b>
0	M	60	1.80
1	F	30	1.60
2	F	10	1.50
3	M	7	1.45
4	M	3	0.73

<i>index</i>	<b>gender</b>	<b>age</b>	<b>height</b>
0	M	60	1.80
3	M	7	1.45
4	M	3	0.73

<i>index</i>	<b>gender</b>	<b>age</b>	<b>height</b>
1	F	30	1.60
2	F	10	1.50

	age	height
gender		
F	20.000000	1.550000
M	23.333333	1.326667

	age		height	
	mean	std	mean	std
gender				
F	20.000000	14.142136	1.550000	0.070711
M	23.333333	31.817186	1.326667	0.545558

`dfg.groupby(by="gender")`

`dfg.groupby(by="gender").mean()`

`dfg.groupby(by="gender").aggregate([np.mean, np.std])`

# Pandas: grouping

*Average expenditures per sector groups*

Use index for grouping:

	2006	2007	2008	2009
Nummer				
B 05-09	0.3	0.5	0.4	
C 10-12	2.5	2.7	2.5	
C 13-15	0.6	0.8	0.8	

inno\_num.T

```
(inno_num.T
 .groupby(by=lambda index: index[0])
 .apply(np.mean))
```

	2006	2007	2008	2009
B	0.300000	0.500000	0.400000	0.200000
C	4.847059	4.988235	5.417647	4.750000
D	1.900000	2.500000	2.300000	2.500000

**10 min break**

# Introduction to Python for data scientists

*Final exercise:*